
EXPLORING MTCP BASED SINGLE-THREADED AND MULTI-THREADED WEB SERVER DESIGN

By

Pijush Chakraborty

Roll: 153050015

Guided By:

Prof. Sriram Srinivasan

Prof. Purushottam Kulkarni



User Space Network Stack

2

□ Why do we need an user-space network stack?

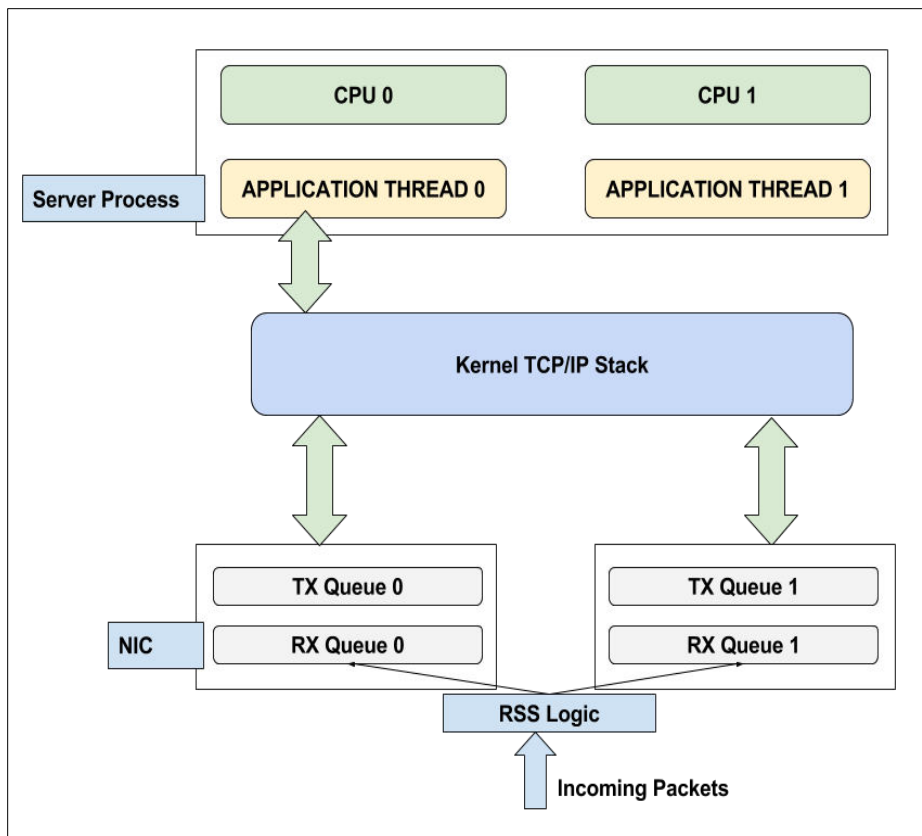


Figure: Kernel Network Stack

- 1 Single Kernel Stack
- 2 System Call Overhead
- 3 Heavy Shared Data-Structures

mTCP: Multicore User-Level TCP Stack

3

- Design of mTCP stack:
 - No Sharing Model
 - Per Core Independent TCP Stacks
 - Per Core Flow Affinity

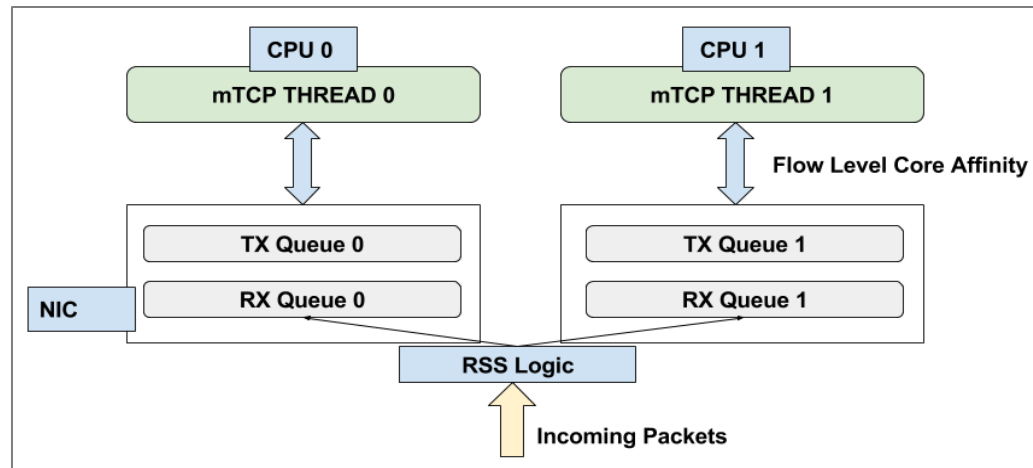
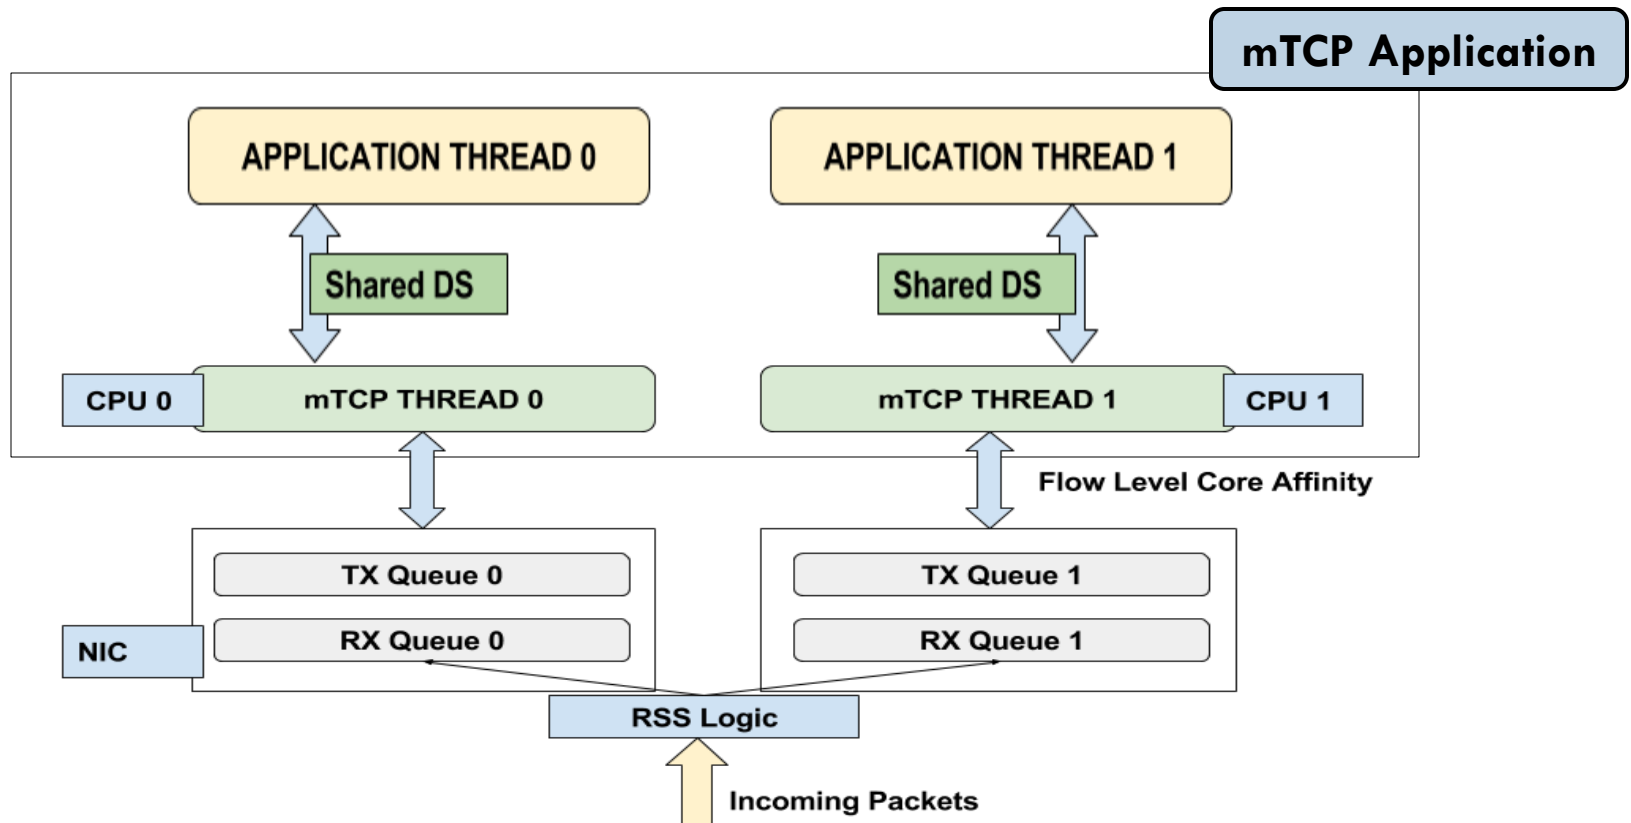


Figure: mTCP Stack

mTCP Based Application Design

4



mTCP Application Design

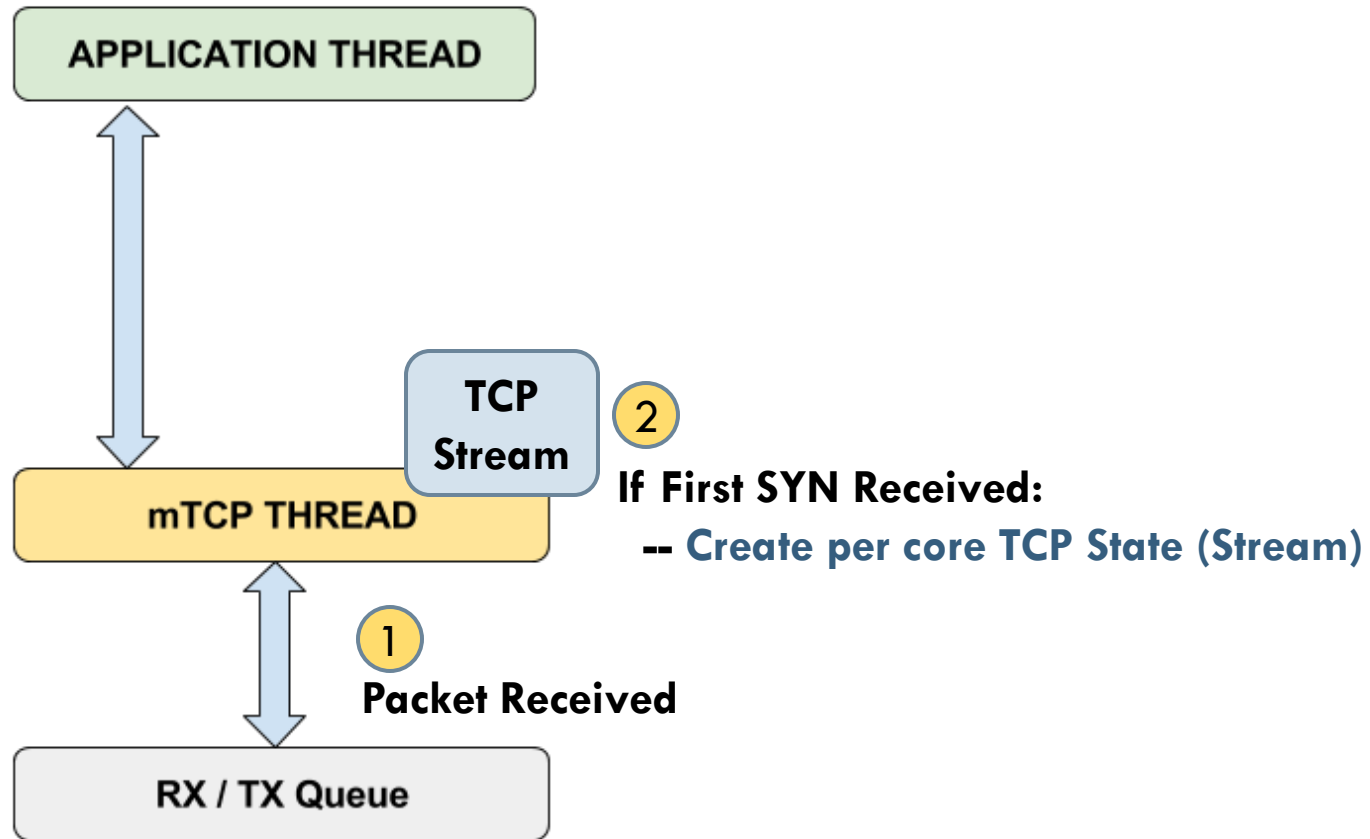
Problem Statement

5

- Compare mTCP based single threaded and multithreaded web server designs
- Answers to find:
 - ▣ Can a master-worker web server design work with current mTCP architecture?
 - ▣ What is the performance overhead in having a shared mTCP architecture?

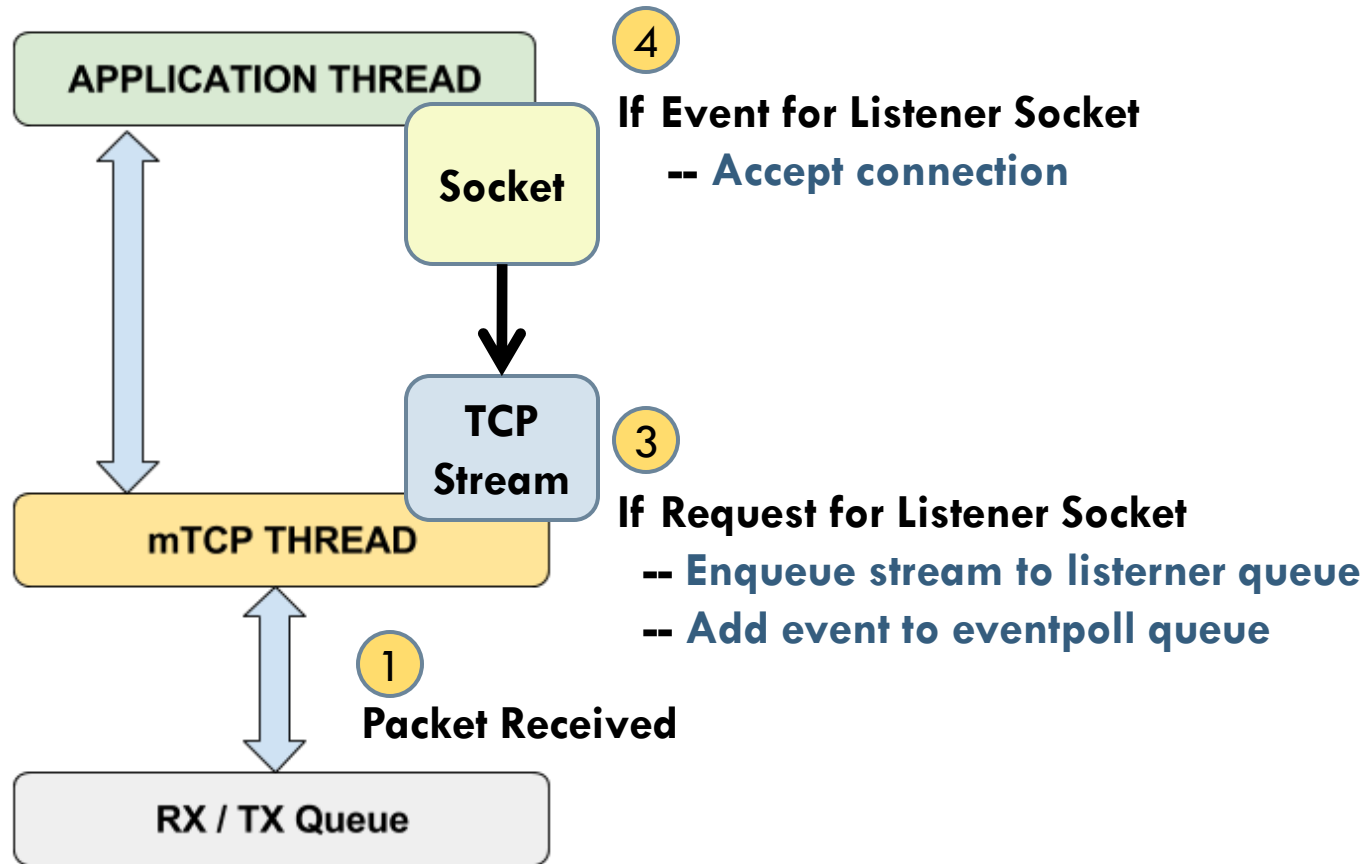
mTCP Single-Threaded Web Server

6



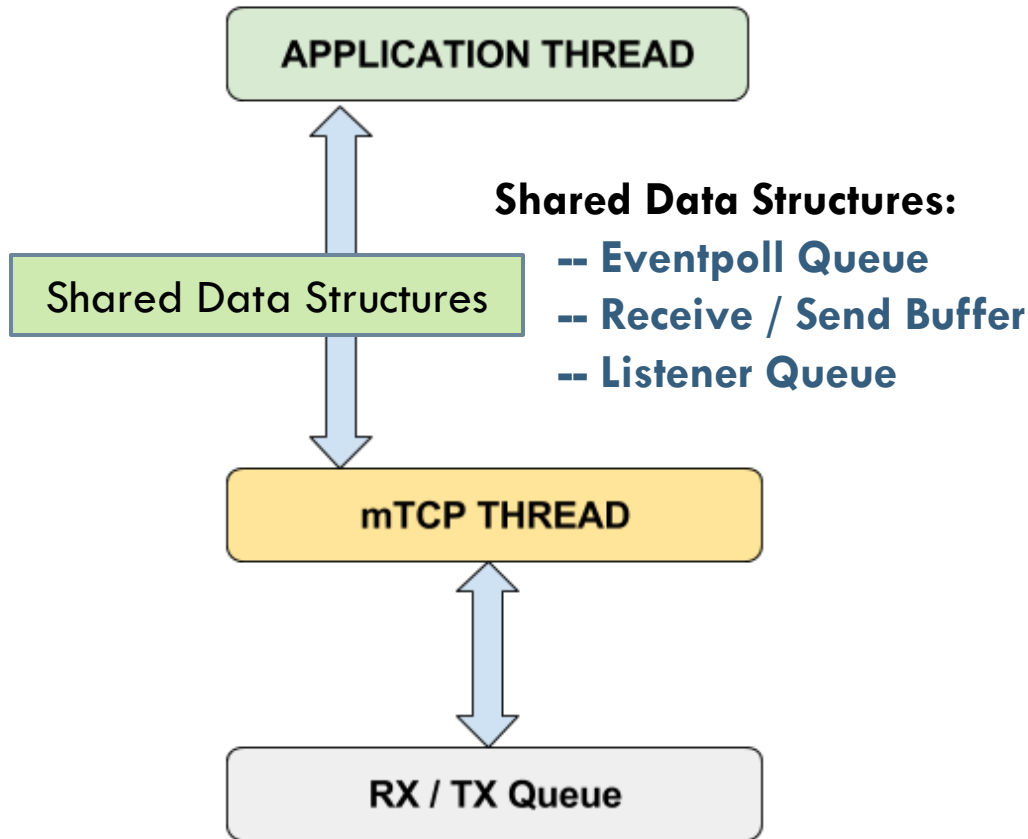
mTCP Single-Threaded Web Server

7



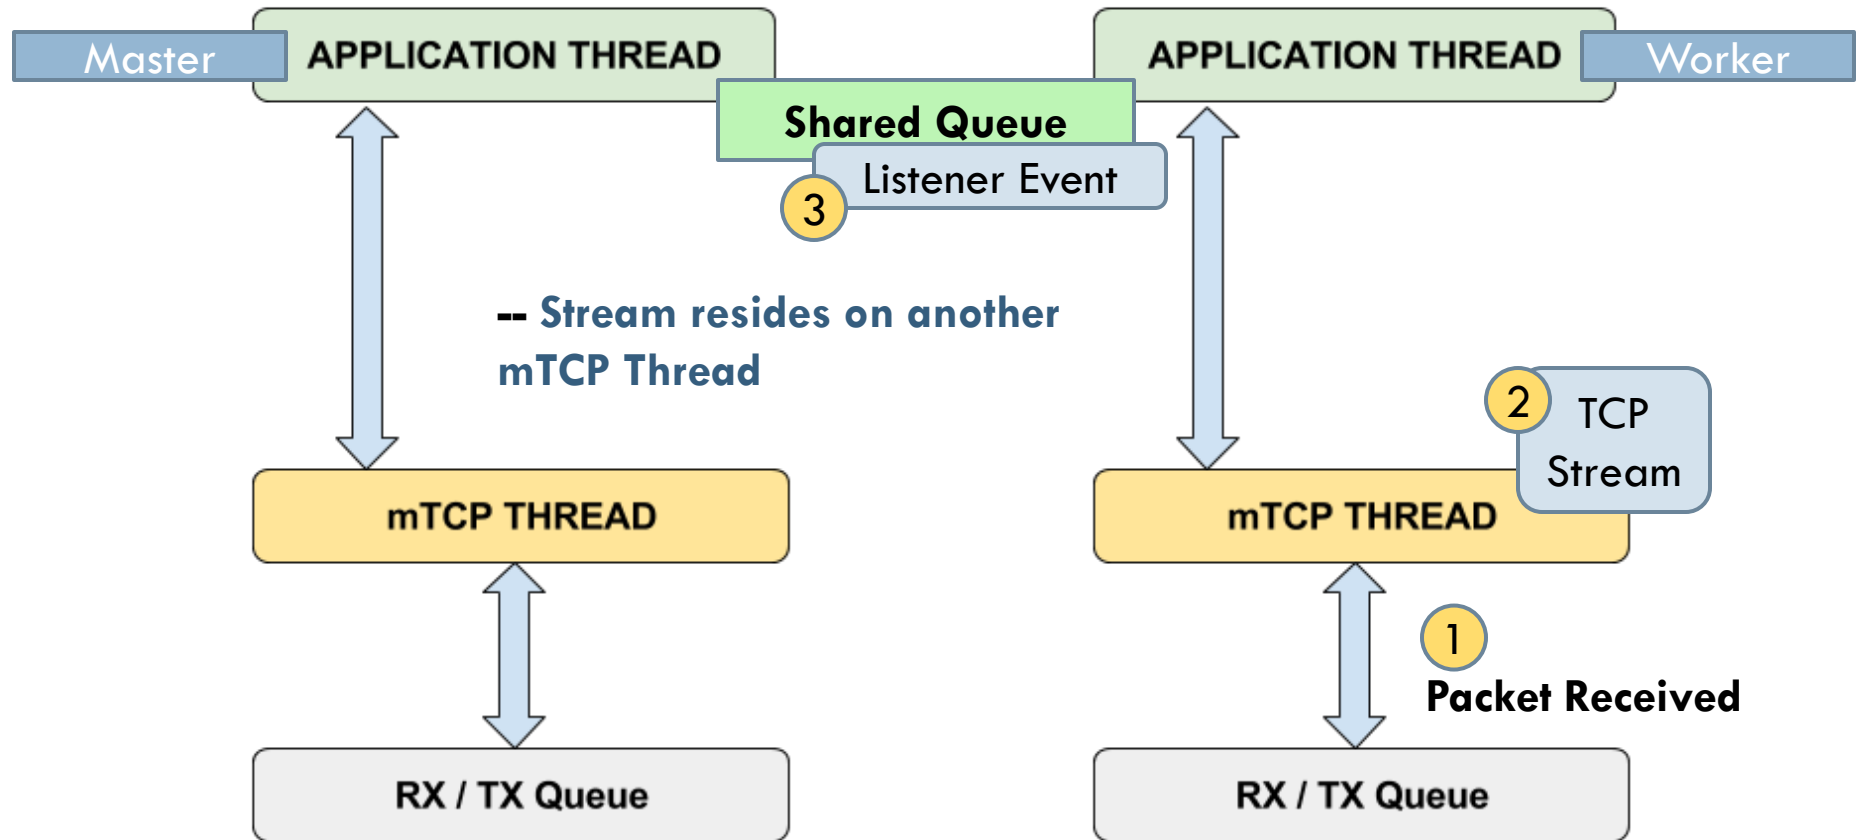
mTCP Single-Threaded Web Server

8



mTCP Multithreaded Web Server

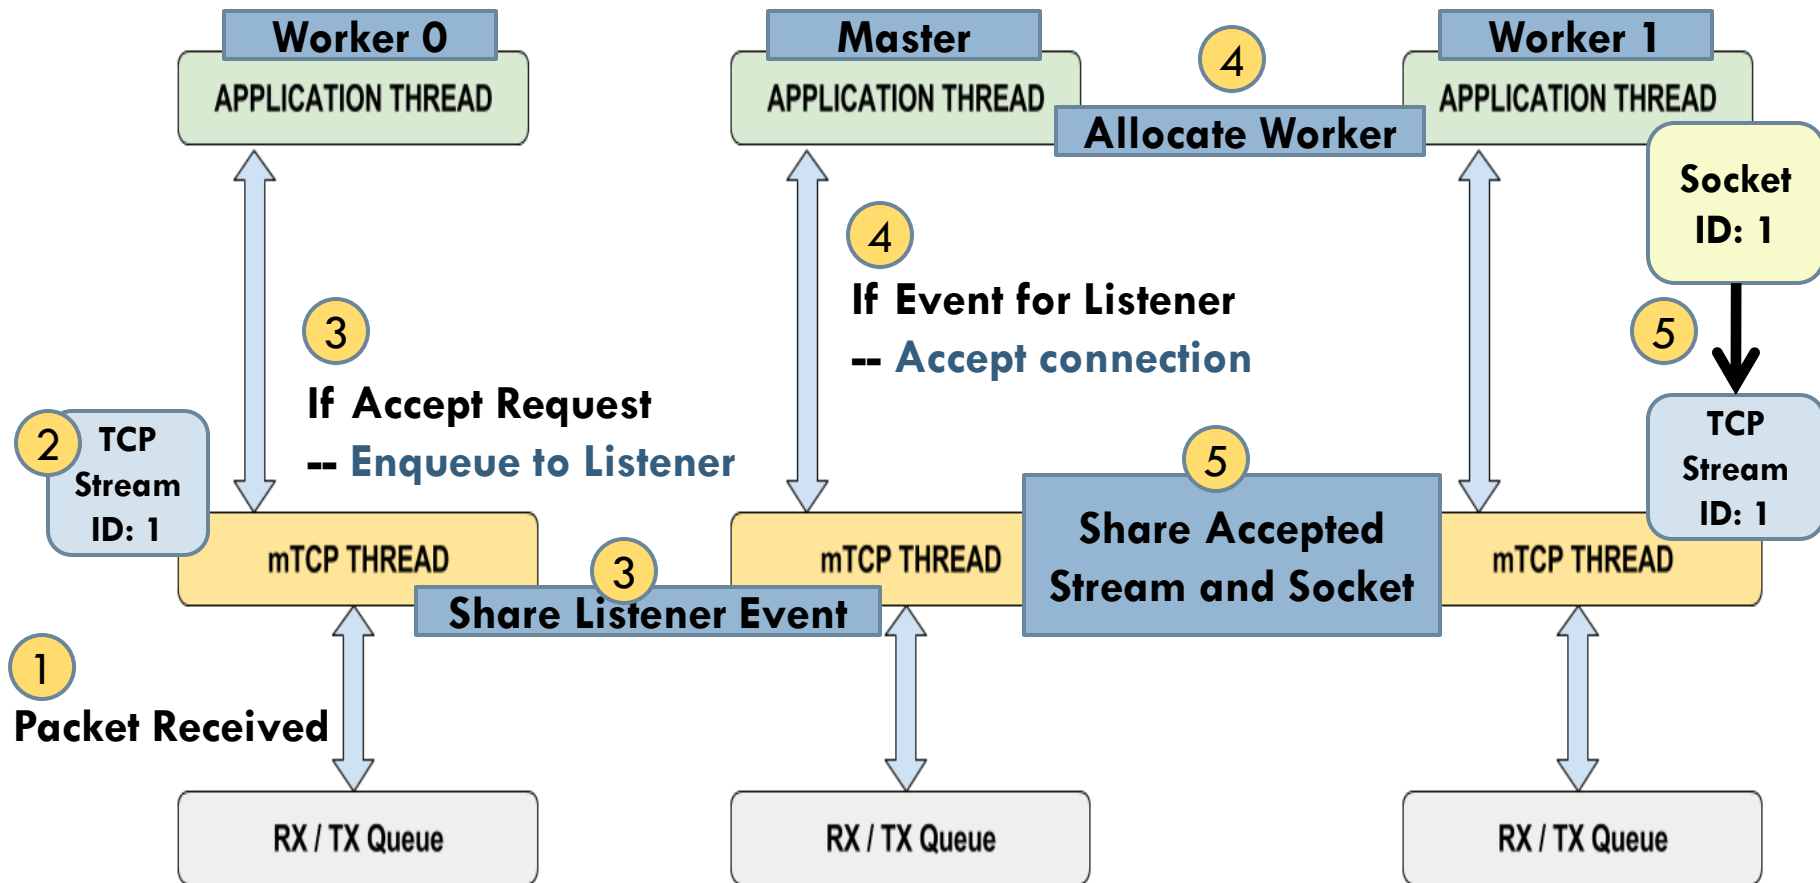
9

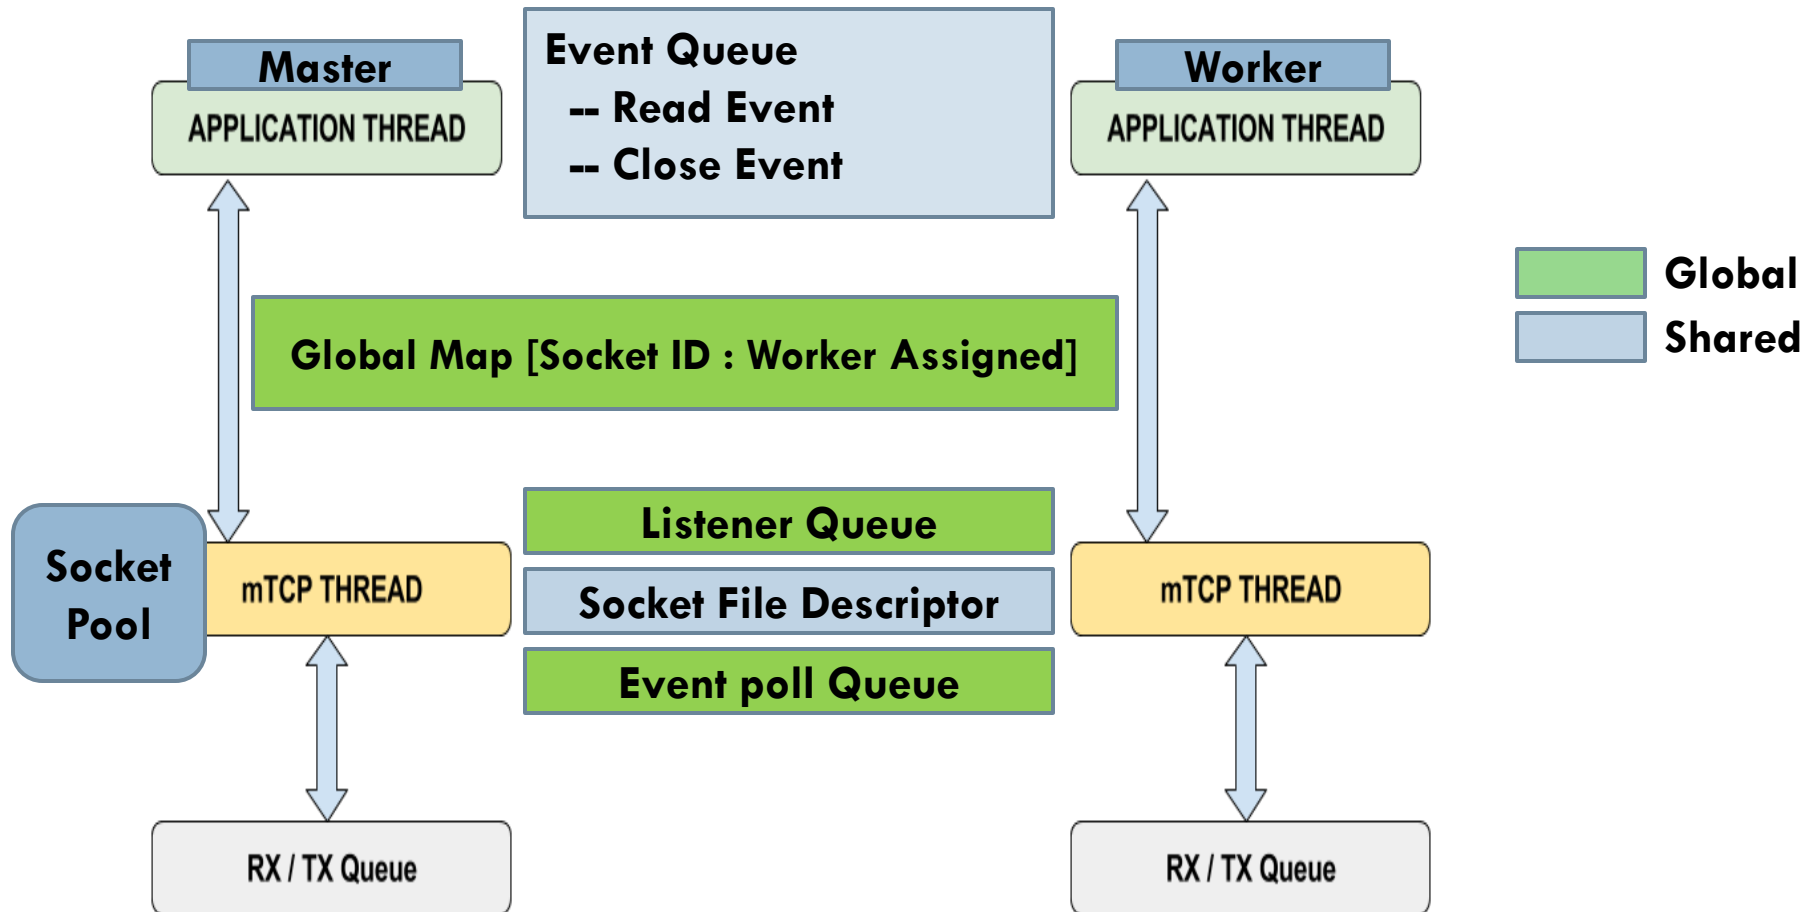


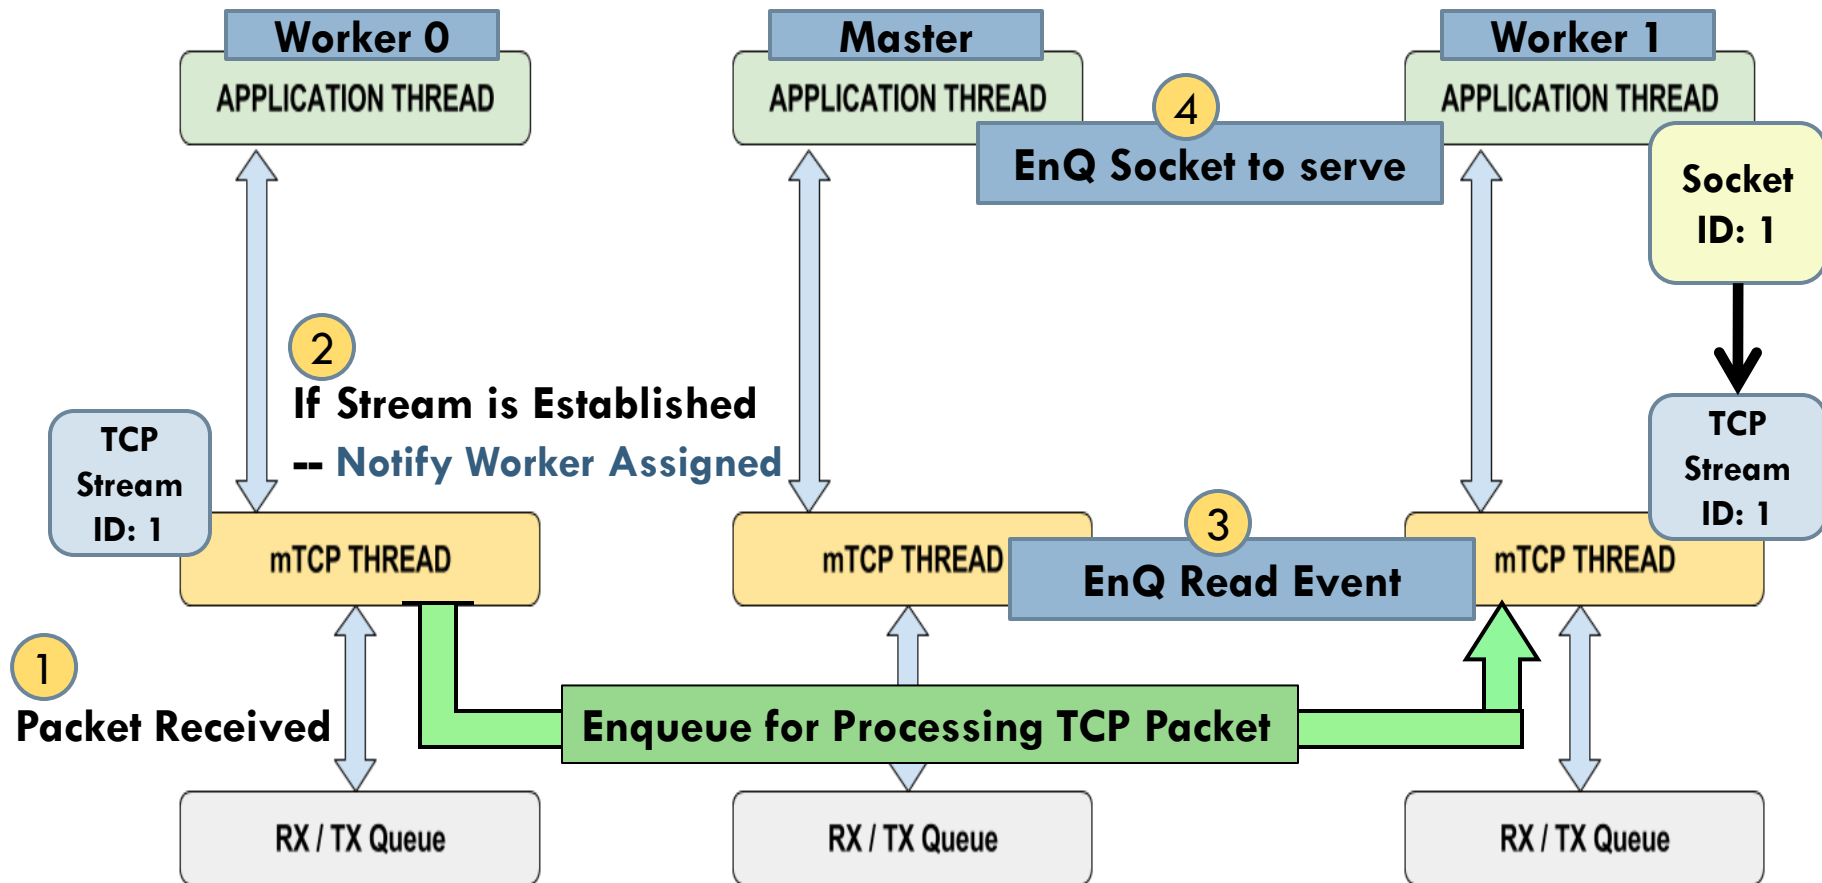
Shared mTCP Network Stack Design

10

- mTCP based Master-Worker designs:
 - ▣ Master controls all events [Read and Accept]
 - Master does the polling and allocates workers
 - ▣ Master controls only Listener Accept Events
 - Master only sets the initial connection
 - Worker works independently after that

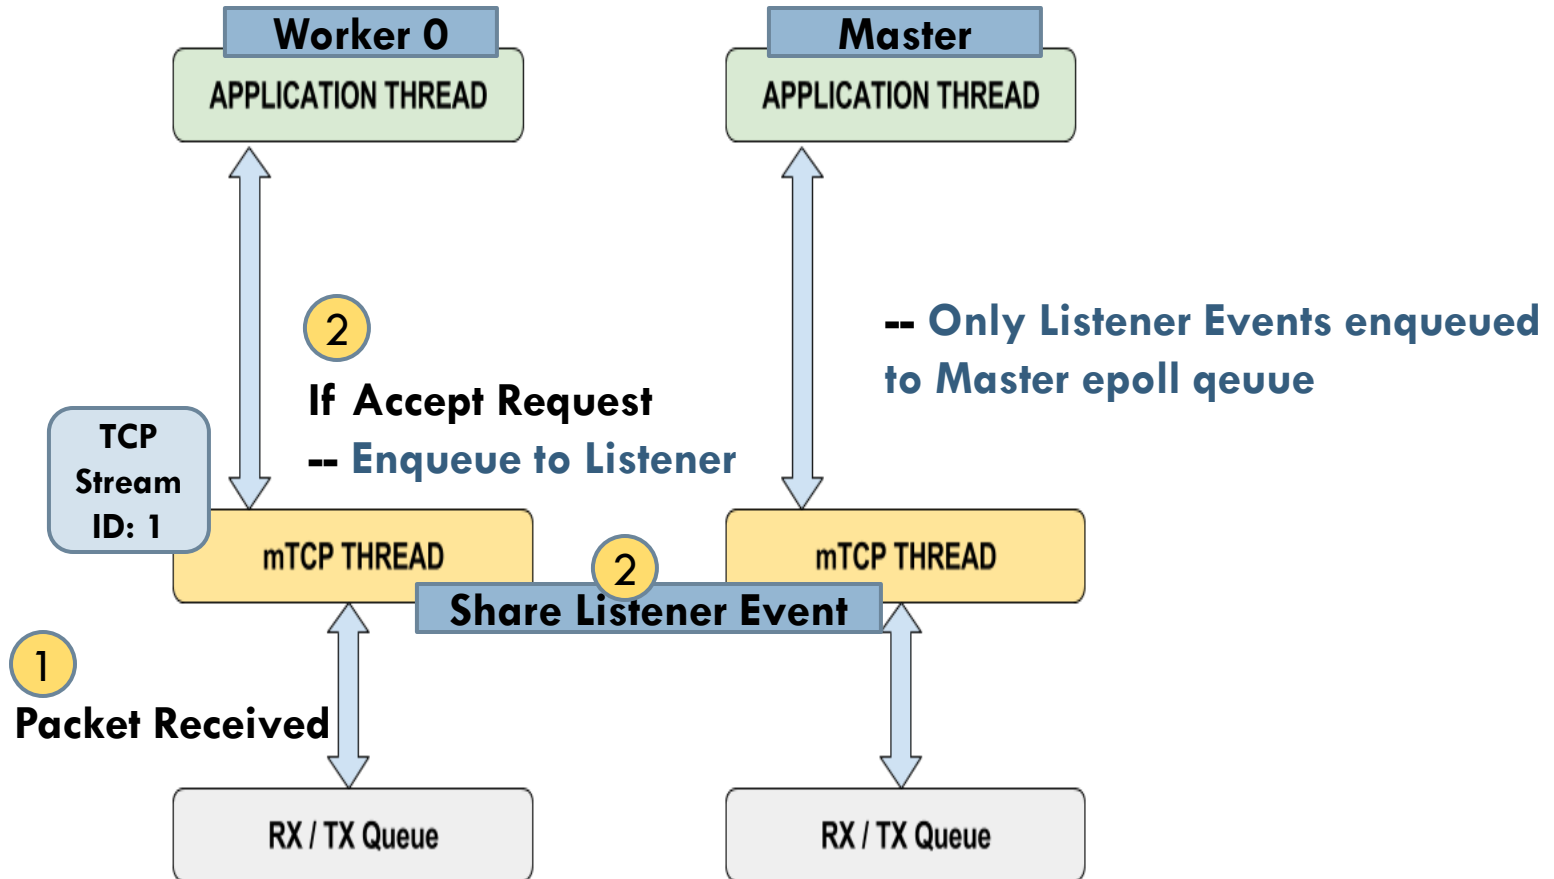






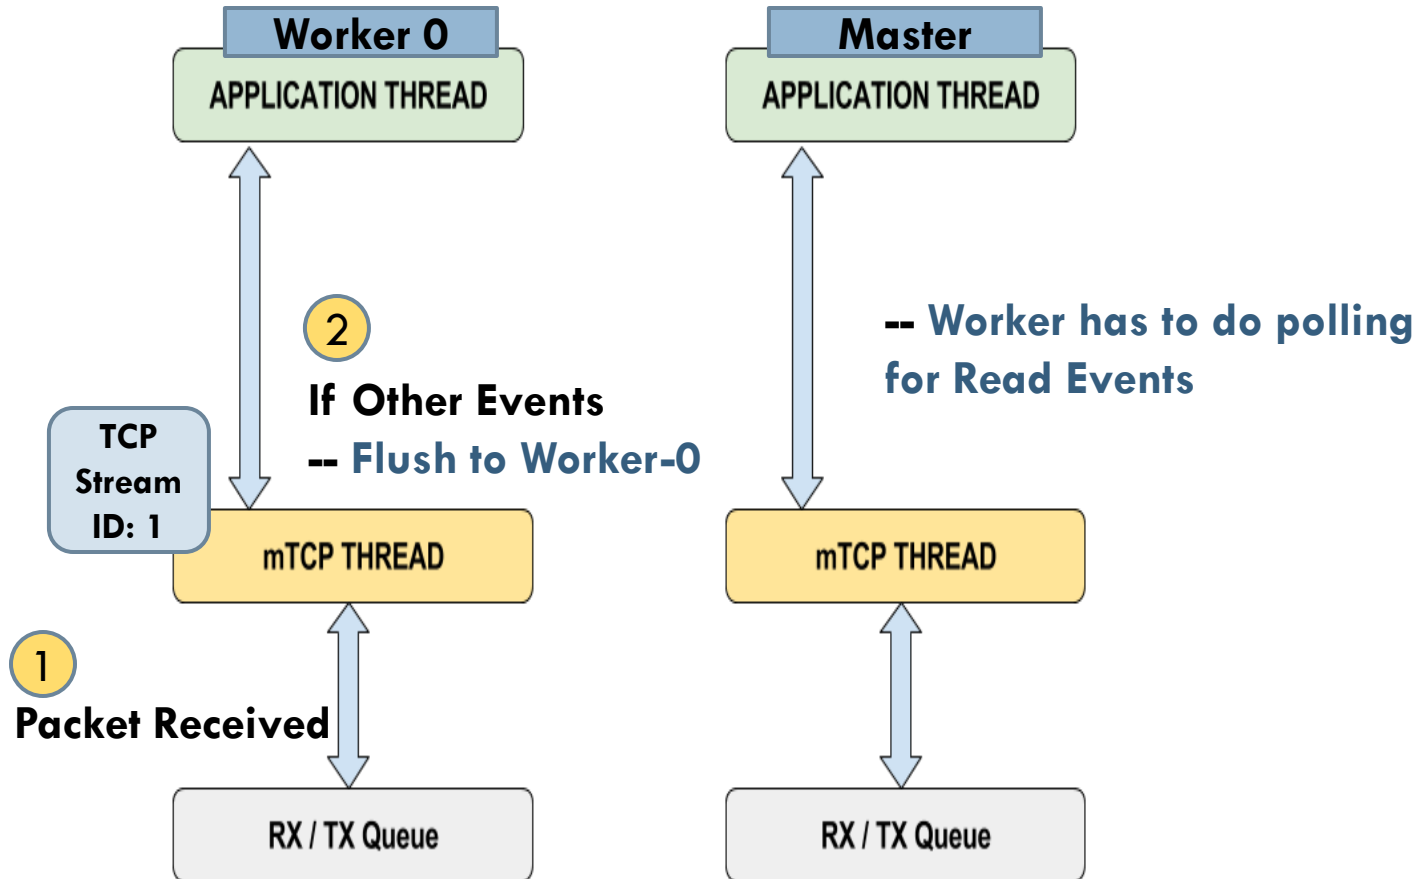
Another Master-Worker Design

14



Another Master-Worker Design

15



Comparing mTCP based Web Servers

16

- Single Threaded Design
 - ▣ Independent server threads working in parallel
 - ▣ Every thread has to do polling
- Multi Threaded Design [Master has Full Control]
 - ▣ Only master thread has to do the epoll
 - ▣ Workers depend on master thread's policies
- Multi Threaded Design [Master Initiates Worker]
 - ▣ Worker too has to do epoll for sockets allocated to it
 - ▣ Worker has full control of connections allocated to it

Experimental Setup

17

- System: 16 Core Xeon Blade Server, 32GB Ram, 1GB NIC
- Experiments to be Done:
 - ▣ Server Throughput based on varying file size requests from client
 - ▣ Server Throughput based on varying no of cores assigned to server


Performance on Varying File Sizes

18

- Client: 2000 open sockets on client side which sends requests of varying sizes for 20 seconds
- Server: The server applications have 7 cores assigned to it

Performance on Varying File Sizes

19



Performance based on Varying File Sizes

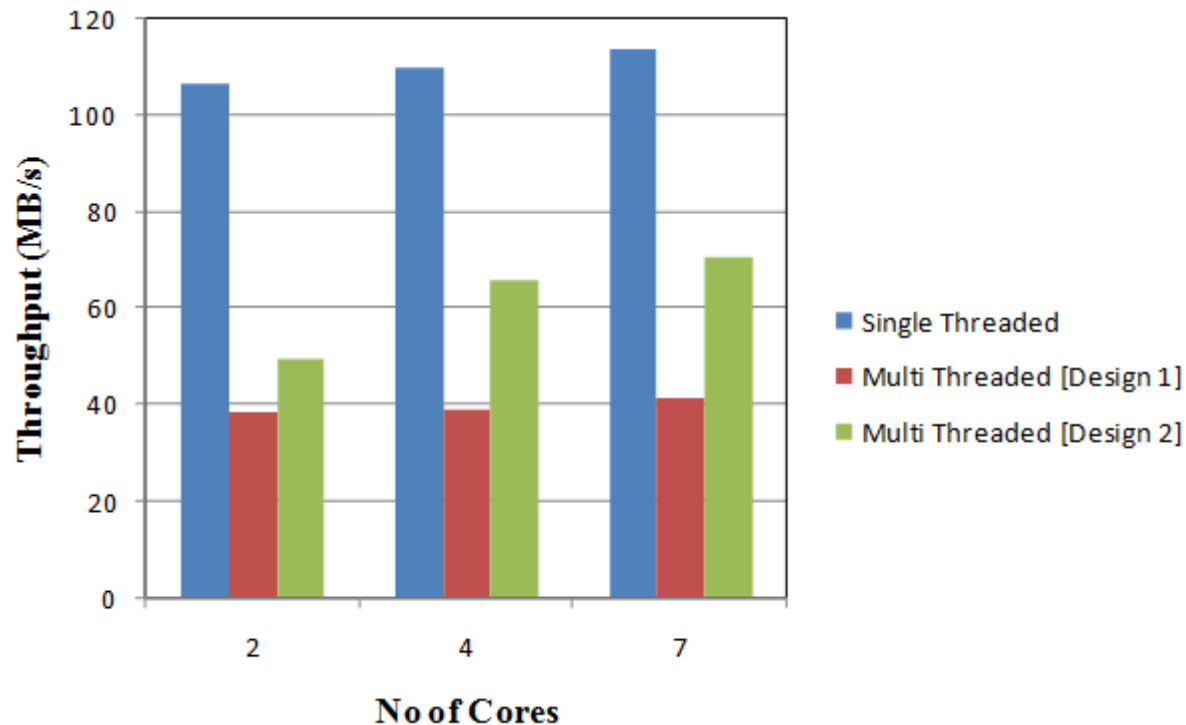
Performance on Varying No. of Cores

20

- Client: 2000 open sockets on client side which sends requests of 1KB file for 20 seconds
- Server: The server application have varying cores assigned to it based on the experiment

Performance on Varying No. of Cores

21



Performance based on varying Number of Cores

Read Copy Update

22

- Never Block Readers
- Writer synchronization left to the developers
- Use of multiple versions of the same data structure

Read Log Update

23

- ❑ Provides a per thread log for concurrent writers
- ❑ Easy to use writer synchronization
- ❑ Provide multiple updates using a single operation

Comparing RCU and RLU

24

- Is RLU better than RCU?
 - ▣ RCU provides better results for Hash-Lists
- Does Node-Size matter in the comparisons?
 - ▣ The relative performance remains the same

Some Other Questions Answered

25

- Do writers need help in RCU protect data structures in the Linux Kernel?
 - ▣ Writer Contention is quite low in the given time frame

- Can Array Based RCU Design help Readers?
 - ▣ For small size objects it performs quite well but the performance drops significantly after the given cache size.

References

- [1] Timothy L. Harris. A pragmatic implementation of non-blocking linked-lists. Proceedings of the 15th International Conference on Distributed Computing, pages 300–314, October 2001.
- [2] Alexander Matveev, Nir Shavit, Pascal Felber, and Patrick Marlier. Read-log-update: a lightweight synchronization mechanism for concurrent programming. 2015.
- [3] Paul E. McKenney. Read-copy update (RCU) usage in Linux kernel. Available: <http://www.rdrop.com/users/paulmck/RCU/linuxusage/rculocktab.html>, October 2006.
- [4] Paul E. McKenney. What is rcu: <https://lwn.net/articles/262464/>. 2007.
- [5] Paul E. McKenney. Rcu linux usage log: <http://www.rdrop.com/paulmck/rcu/linuxusage/linux-4.3.rcua>. 2015.
- [6] Paul E. McKenney. Read-mostly research in 2015: <https://lwn.net/articles/667593>. 2015.
- [7] Paul E. McKenney. Some more details on read-log-update: <https://lwn.net/articles/667720>. 2015.
- [8] Paul E. McKenney and John D. Slingwine. Read-copy update: Using execution history to solve concurrency problems. In Parallel and Distributed Computing and Systems, pages 509–518, Las Vegas, NV, October 1998.
- [9] Maged M. Michael. Hazard pointers: Safe memory reclamation for lock-free objects. IEEE Transactions on Parallel and Distributed Systems, 15(6):491–504, June 2004.
- [10] John D. Valois. Lock-free linked lists using compare-and-swap. , Proceedings of the 14th annual ACM symposium on Principles of distributed computing, pages 214–222, August 1995.

Questions?