
TOWARDS PERFORMANCE ANALYSIS OF NON-BLOCKING CONCURRENT DATA STRUCTURES IN THE LINUX KERNEL

By

Pijush Chakraborty

Roll: 153050015

Guided By:

Prof. Sriram Srinivasan

Prof. Purushottam Kulkarni



Contents



- Introduction
- Problem Statement
- Establishing Claims
- Conclusion and Stage-2 Plans

Lock Free Data Structures


- Allows multiple threads to work in parallel on a particular data structure of interest.
- No wastage of CPU-cycles for waiting for a mutex lock.
- Other major problems to consider:
 - ▣ ABA problem
 - ▣ Memory Reclamation

Read Mostly Lock Free Data Structures

- Lock Free Data Structures are often compiled with memory reclamation methods.
- Research for Read Mostly Lock Free Data Structures
 - ▣ Read Copy Update Mechanism
 - ▣ Read Log Update Mechanism


Read Copy Update

- Never Block Readers
- Writer synchronization left to the developers
- The Mechanism:
 - ▣ Read:




Read Copy Update

- Never Block Readers
- Writer synchronization left to the developers
- The Mechanism:
 - Copy:



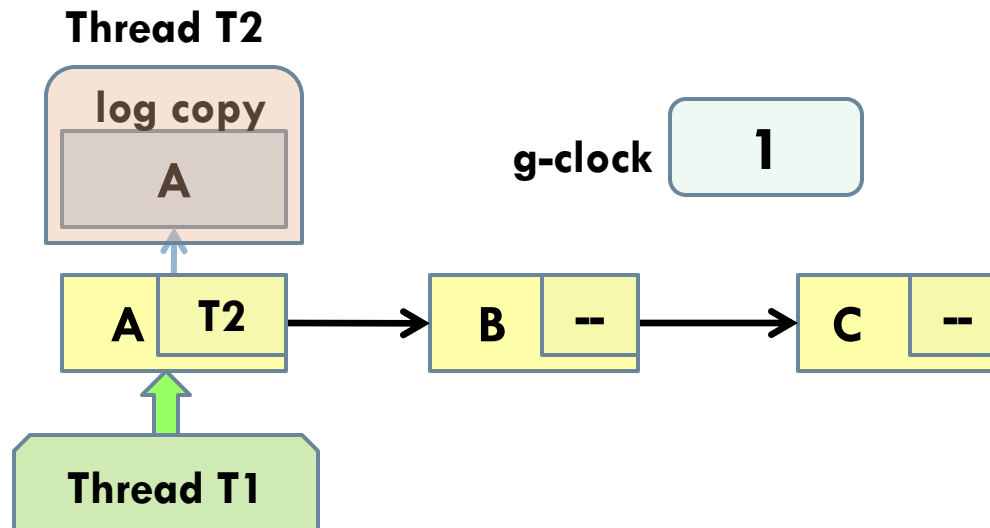
Read Copy Update

- Never Block Readers
- Writer synchronization left to the developers
- The Mechanism:
 - ▣ Update:



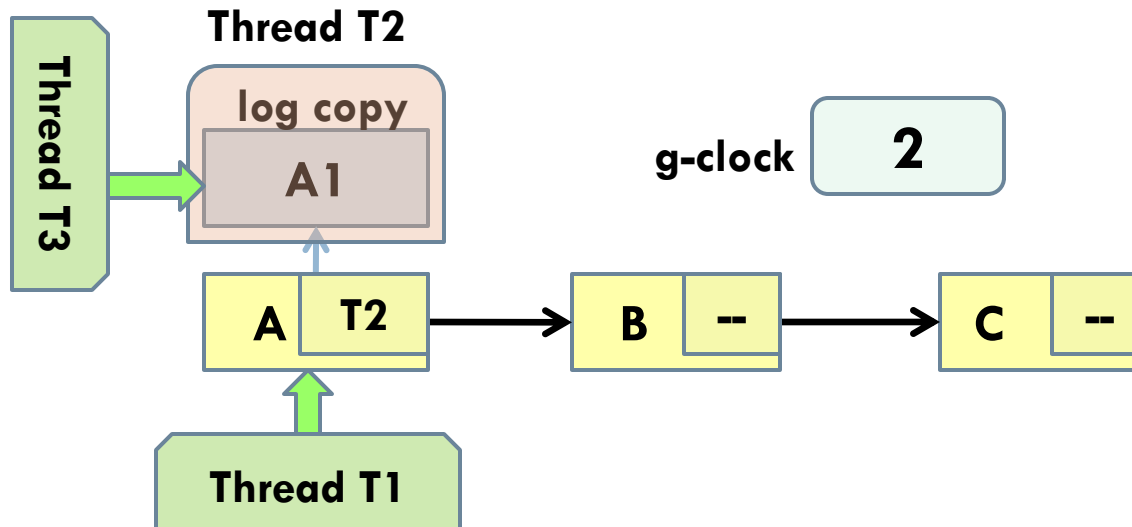
Read Log Update

- Provides a per thread log for concurrent writers
- Easy to use writer synchronization
- The Mechanism:
 - ▣ Read-Log:



Read Log Update

- Provides a per thread log for concurrent writers
- Easy to use writer synchronization
- The Mechanism:
 - ▣ Update:



Contents



- Introduction
- **Problem Statement**
- Establishing Claims
- Conclusion and Stage-2 Plans

Problem Statement

- Compare existing Read Mostly Lock Free Mechanisms.
- Answers to find:
 - ▣ Should writers be helped to improve overall performance?
 - ▣ Can we improve existing list based semantics in use to improve cache utilization?


Contents

- Introduction
- Problem Statement
- **Establishing Claims**
- Conclusion and Stage-2 Plans


Verifying Read-Log-Update Claims

- Questions:
 - ▣ Is RLU better than RCU?
 - ▣ Does Node-Size matter in the comparisons?
- Experiments for Comparing RCU and RLU
 - ▣ System: 16 core blade server(Intel Xeon) supporting 16 Hardware threads.
 - ▣ Compared using benchmark used by RLU authors

Comparing Linked Lists




Comparison with no updates




Comparison with 40% updates

Comparing Hash Lists


- Hash Lists protected by RCU and RLU in use



Comparing Hash Lists




Comparison with no updates



Comparison with 40% updates


Does Node Size Matter?

Comparison between RCU and RLU



Linked List with 20% updates

Comparison between RCU and RLU



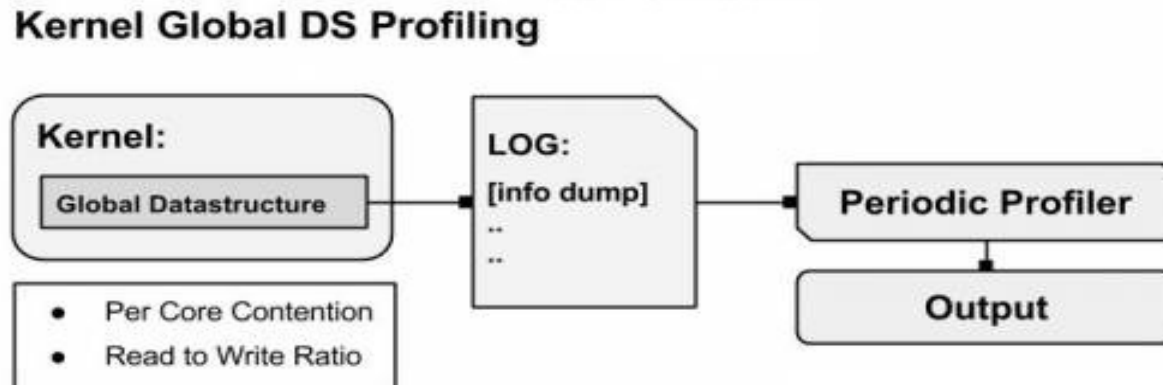
Linked List with 40% updates

Kernel Side of the Story

- Question:
 - ▣ Do the writers need any optimization?
- Understand the usage behavior of the currently used RCU protected data structures
 - ▣ Read-Write Ratio
 - ▣ Write Contention

Experiments:

- System: 4 core Intel i5 processor(1.2GHz), 8GB Ram
- Design:





- Workload:
 - ▣ Web Server Workload
 - ▣ Process List Stress Workload

Profiling Data Structures

- List to Track: epoll file descriptor list
- Server in use:
 - ▣ Nginx local server that uses epoll to serve multiple client requests
- Client Statistics:
 - ▣ 1000 parallel clients created every 5 seconds
 - ▣ Each client making 10000 requests in parallel to localhost.

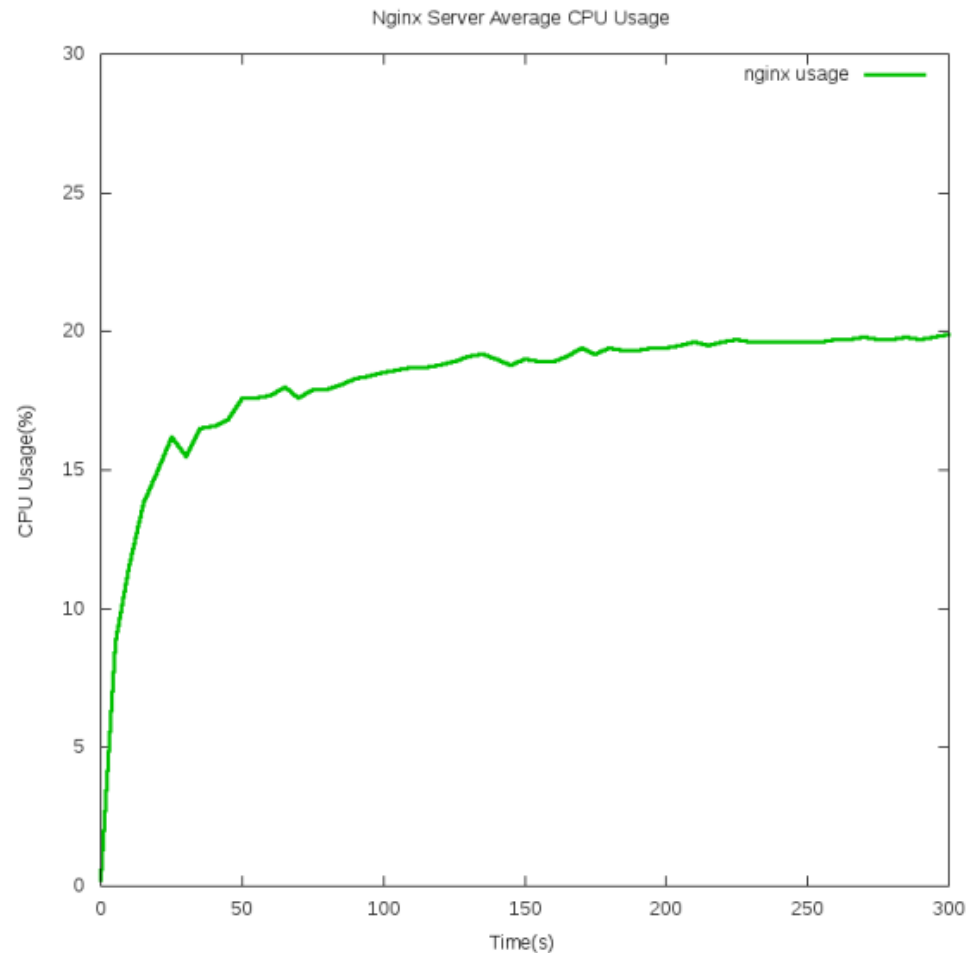
Profiling Data Structures

Workload Load: Web Server Workload



Profiling Data Structures

□ Workload Load: Web Server Workload


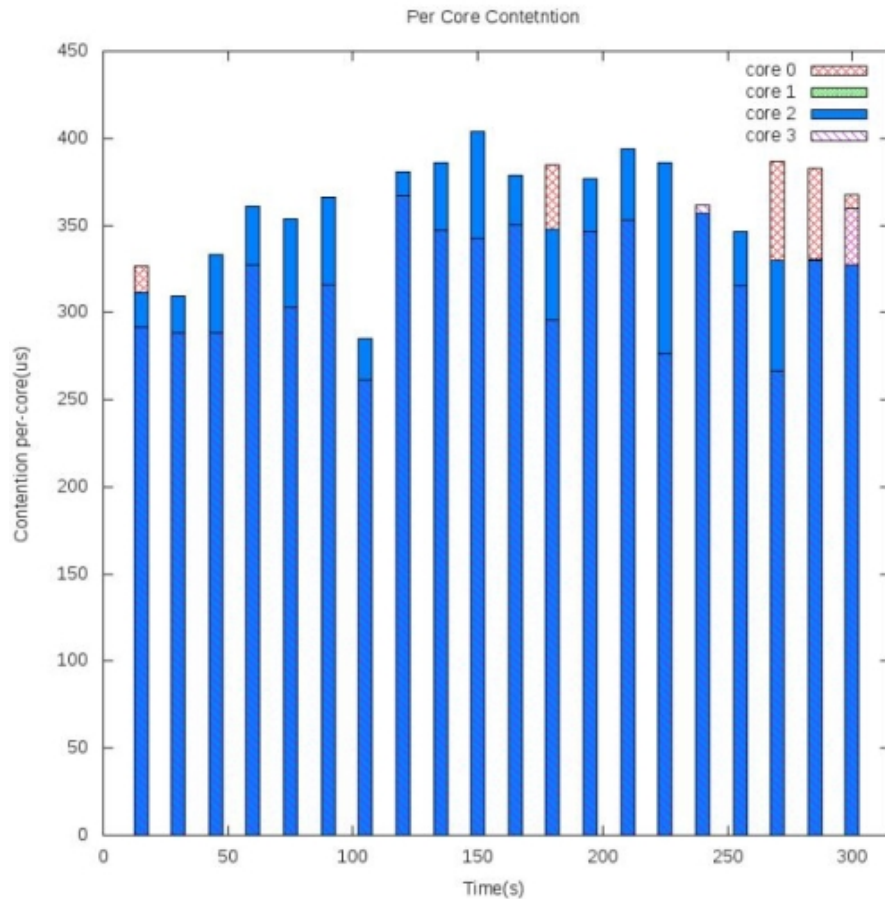


Profiling Data Structures

- List to Track: List of all processes
- Synthetic Workload:
 - ▣ 250 parallel threads
 - ▣ Each thread creating 500 empty threads every 5 seconds

Profiling Data Structures

Workload Load: Synthetic Stress Workload



Helping the Readers

- Questions:
 - ▣ Can we provide a cache friendly way of traversal?
 - ▣ Should the list based semantics be changed?
- Kernel data structures uses list based semantics
- But, arrays provide better cache utilization than lists

Observations

- Are the lists ordered?

Name of List Head	Insertion Behavior	Deletion Behavior
List of process with same thread-id	Inserts at head	Delete anywhere
List for epoll file descriptors	Inserts at tail	Delete anywhere
List of all running processes	Inserts at head	Delete anywhere

Table 1: Usage Behavior Observations

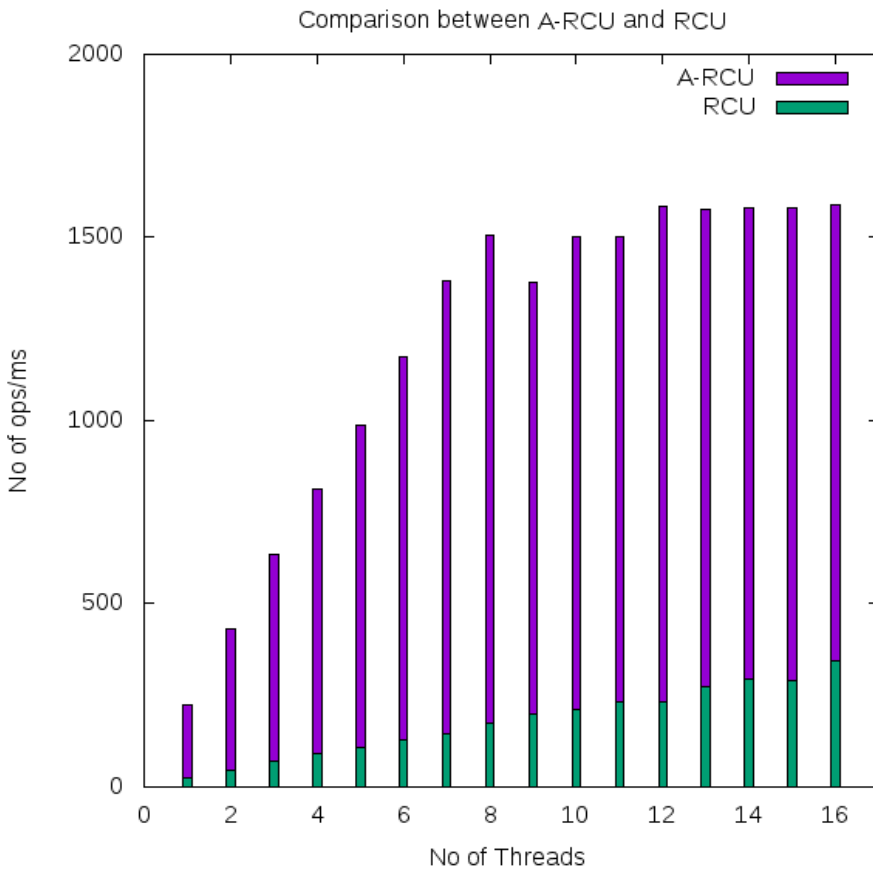
Usage Statistics	Linked List	Hash List
No of such RCU protected data structures in the Kernel	236	71
Inserts at list tail	109	–
Inserts at head	127	71

Table 2: Usage Statistics Assumptions

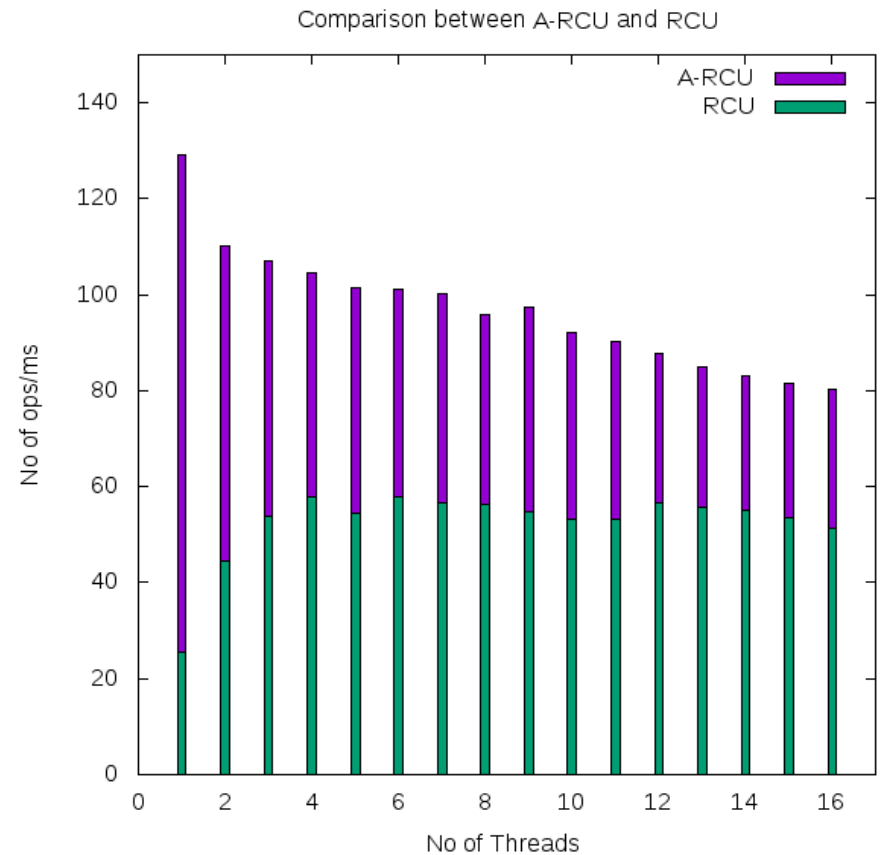
RCU Protected Arrays

- Can RCU protected Arrays prove better?
- Experiment:
 - ▣ Comparison using RCU array(A-RCU) and RCU protected linked list
 - ▣ Benchmark used by RLU authors for comparing RCU and RLU
 - ▣ System: 16 core Xeon Blade Server processor supporting 16 hardware threads.

RCU Protected Arrays

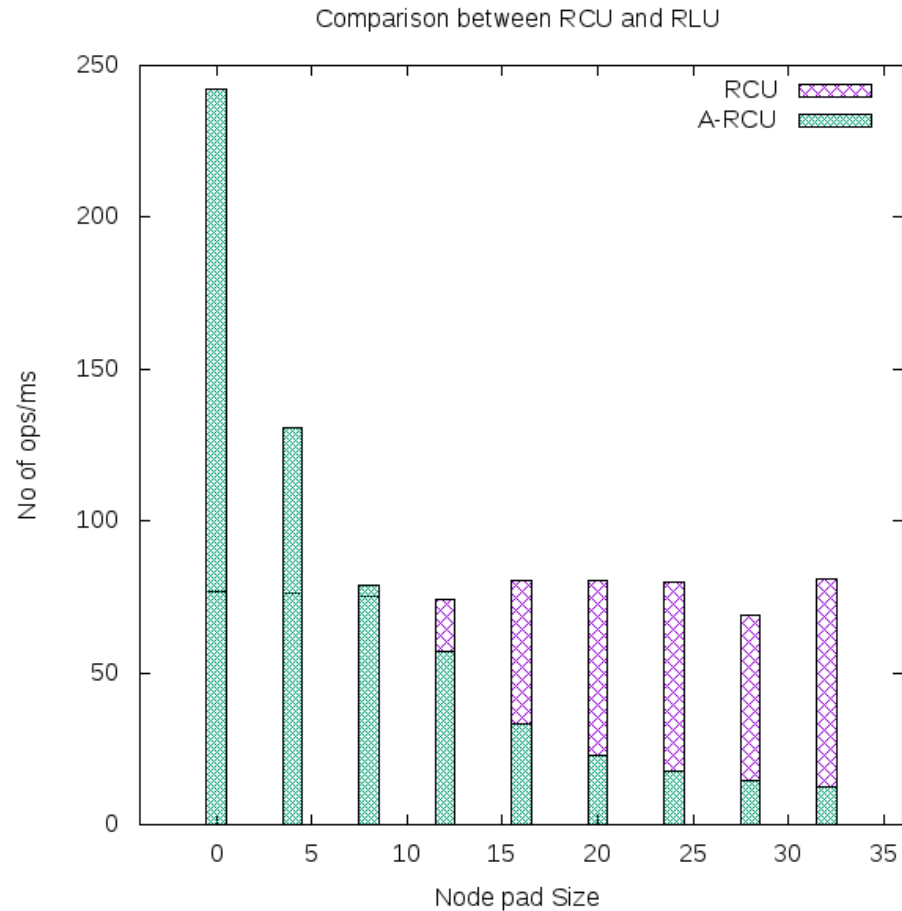


Comparison with no updates



Comparison with 40% updates

Does Node Size Matter?



Comparison with 20% updates

Conclusion

- ❑ Studied the behavioral aspects of RCU Usage in the Linux Kernel.
- ❑ Need to look at Reader Perspective in Stage-2.
- ❑ Design a lock free mechanism from that supports:
 - ▣ More cache utilization
 - ▣ Supports larger node size
- ❑ Change a particular Kernel subsystem to use the array based design and compare with existing model.

References

- [1] Timothy L. Harris. A pragmatic implementation of non-blocking linked-lists. Proceedings of the 15th International Conference on Distributed Computing, pages 300–314, October 2001.
- [2] Alexander Matveev, Nir Shavit, Pascal Felber, and Patrick Marlier. Read-log-update: a lightweight synchronization mechanism for concurrent programming. 2015.
- [3] Paul E. McKenney. Read-copy update (RCU) usage in Linux kernel. Available: <http://www.rdrop.com/users/paulmck/RCU/linuxusage/rculocktab.html>, October 2006.
- [4] Paul E. McKenney. What is rcu: <https://lwn.net/articles/262464/>. 2007.
- [5] Paul E. McKenney. Rcu linux usage log: <http://www.rdrop.com/paulmck/rcu/linuxusage/linux-4.3.rcua>. 2015.
- [6] Paul E. McKenney. Read-mostly research in 2015: <https://lwn.net/articles/667593>. 2015.
- [7] Paul E. McKenney. Some more details on read-log-update: <https://lwn.net/articles/667720>. 2015.
- [8] Paul E. McKenney and John D. Slingwine. Read-copy update: Using execution history to solve concurrency problems. In Parallel and Distributed Computing and Systems, pages 509–518, Las Vegas, NV, October 1998.
- [9] Maged M. Michael. Hazard pointers: Safe memory reclamation for lock-free objects. IEEE Transactions on Parallel and Distributed Systems, 15(6):491–504, June 2004.
- [10] John D. Valois. Lock-free linked lists using compare-and-swap. , Proceedings of the 14th annual ACM symposium on Principles of distributed computing, pages 214–222, August 1995.